

Conectando ao MySQL e listando as bases.

```
(base) [hadoop@dataserver ~]$ sqoop list-databases --connect jdbc:mysql://localhost:3306/ --username root -P
2020-07-14 09:30:36,065 INFO sqoop.Sqoop: Running Sqoop version: 1.4.7
Enter password:
2020-07-14 09:30:42,636 INFO manager.MySQLManager: Preparing to use a MySQL streaming resultset.
Loading class 'com.mysql.jdbc.Driver'. This is deprecated. The new driver class is 'com.mysql.cj.jdbc.Driver'. The driver is automatically registered
SPI and manual loading of the driver class is generally unnecessary.
mysql
information_schema
performance_schema
sys
testedb
adventureworks
```

Listando as tabelas do banco.

```
(base) [hadoop@dataserver ~]$ sqoop list-tables --connect jdbc:mysql://localhost:3306/adventureworks --username root -P
2020-07-14 11:01:58,775 INFO sqoop.Sqoop: Running Sqoop version: 1.4.7
Enter password:
2020-07-14 11:02:03,394 INFO manager.MySQLManager: Preparing to use a MySQL streaming resultset.
Loading class 'com.mysql.jdbc.Driver'. This is deprecated. The new driver class is 'com.mysql.cj.jdbc.Driver'. The driver is automatically registered via the
SPI and manual loading of the driver class is generally unnecessary.
address
address_type
awbuildversion
billofmaterials
contact
contactcreditcard
contacttype
countryregion
...
specialofferproduct
stateprovince
store
storecontact
transactionhistory
transactionhistoryarchive
unitmeasure
vendor
vendoraddress
vendorcontact
workorder
workorderrouting
(base) [hadoop@dataserver ~]$
```

Importando todas as tabelas para o HDFS.

```
(base) [hadoop@dataserver ~]$ sqoop import-all-tables --connect jdbc:mysql://localhost:3306/adventureworks --username root -P -m 1
2020-07-14 16:16:50,136 INFO sqoop.Sqoop: Running Sqoop version: 1.4.7
Enter password:
2020-07-14 16:16:54,215 INFO manager.MySQLManager: Preparing to use a MySQL streaming resultset.
Loading class 'com.mysql.jdbc.Driver'. This is deprecated. The new driver class is 'com.mysql.cj.jdbc.Driver'. The driver is automatically registered via the
SPI and manual loading of the driver class is generally unnecessary.
2020-07-14 16:16:54,680 INFO tool.CodeGenTool: Beginning code generation
2020-07-14 16:16:54,691 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM 'address' AS t LIMIT 1
2020-07-14 16:16:54,711 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM 'address' AS t LIMIT 1
2020-07-14 16:16:54,720 INFO orm.CompilationManager: HADOOP_MAPRED_HOME is /opt/hadoop
Note: /tmp/sqoop-hadoop/compile/8f7eb382d17e9d61cc2806ecb9f11c37/address.java uses or overrides a deprecated API.
Note: Recompile with -Xlint:deprecation for details.
2020-07-14 16:16:55,559 INFO orm.CompilationManager: Writing jar file: /tmp/sqoop-hadoop/compile/8f7eb382d17e9d61cc2806ecb9f11c37/address.jar
2020-07-14 16:16:55,567 WARN manager.MySQLManager: It looks like you are importing from mysql.
2020-07-14 16:16:55,567 WARN manager.MySQLManager: This transfer can be faster! Use the --direct
2020-07-14 16:16:55,567 WARN manager.MySQLManager: option to exercise a MySQL-specific fast path.
2020-07-14 16:16:55,568 INFO manager.MySQLManager: Setting zero DATETIME behavior to convertToNull (mysql)
2020-07-14 16:16:55,572 INFO mapreduce.ImportJobBase: Beginning import of address
2020-07-14 16:16:55,573 INFO Configuration.deprecation: mapred.job.tracker is deprecated. Instead, use mapreduce.job.tracker.address
2020-07-14 16:16:55,682 INFO Configuration.deprecation: mapred.jar is deprecated. Instead, use mapreduce.job.jar
```

Job MapReduce.

```
FILE: Number of write operations=0
HDFS: Number of bytes read=87
HDFS: Number of bytes written=9603681
HDFS: Number of read operations=6
HDFS: Number of large read operations=0
HDFS: Number of write operations=2
HDFS: Number of bytes read erasure-coded=0

Job Counters
Launched map tasks=1
Other local map tasks=1
Total time spent by all maps in occupied slots (ms)=14853
Total time spent by all reduces in occupied slots (ms)=0
Total time spent by all map tasks (ms)=14853
Total vcore-milliseconds taken by all map tasks=14853
Total megabyte-milliseconds taken by all map tasks=15209472

Map-Reduce Framework
Map input records=67131
Map output records=67131
Input split bytes=87
Spilled Records=0
Failed Shuffles=0
Merged Map outputs=0
GC time elapsed (ms)=62
CPU time spent (ms)=2920
Physical memory (bytes) snapshot=377614336
Virtual memory (bytes) snapshot=2602094592
Total committed heap usage (bytes)=308281344
Peak Map Physical memory (bytes)=377614336
Peak Map Virtual memory (bytes)=2602094592

File Input Format Counters
Bytes Read=0
File Output Format Counters
Bytes Written=9603681
2020-07-14 18:00:35,715 INFO mapreduce.ImportJobBase: Transferred 9.1588 MB in 88.2726 seconds (106.2458 KB/sec)
2020-07-14 18:00:35,716 INFO mapreduce.ImportJobBase: Retrieved 67131 records.
(base) [hadoop@dataserver ~]$
```